

# 1 How to get LSV clusters running - Step by Step Guide

2

## 3 Prerequisites:

4 - OS: any Linux distribution, or alternatively WSL (Windows Subsystem for Linux) if you're currently on Windows (how to WSL: <https://learn.microsoft.com/en-us/windows/wsl/install> )

5 (ssh should also be working out of the box on the Windows command line, but I don't know if there are any intricacies with that and if it will work just as good as on Linux)

6

7 - very basic knowledge of how to use the Linux command line (necessary to connect to (using ssh) and work on the cluster nodes, which run on Ubuntu)

8 (I will walk you through the basics as good as I can)

9

10 - a copy of the project files for the cluster that I put here: [LINK HERE]

11

12 - LSV account, working and password changed from master to your own password

13

14 - ssh installed (should come preinstalled on pretty much any Linux disto, but if not, install via the command 'apt install ssh' (e.g. for Ubuntu)

15 or use any other package manager of your choice)

16

17 - note: wherever I have put the [yourUsername] placeholder, replace it with just your username, drop the square brackets (and don't put any whitespaces in or something like that)

18

19

## 20 Disclaimer:

21

22 I have explicitly written this guide assuming you don't know how to operate a Linux terminal, to make sure everyone understands what we are doing, even without having any Linux experience.

23 I tried to go into great detail on each step, explaining the commands you need and how to use them, even for the most basic commands.

24 If you already are using Linux and have some experience working with the Linux terminal, much of this guide will feel redundant and overly detailed.

25 I made this decision because I wanted to make sure that everyone can not just reproduce but also understand what they are doing, even if they have no Linux experience.

26 (And because I remember having asked about operating systems and being told that we have 2 or 3 persons on Windows in the group.)

27

28 Please also check out the lsv\_cluster\_nodes\_structure.png infographic I made (in the sopro25 folder), which visualizes the remote LSV nodes we will be using and connecting to.

29

30 This guide is the compiled documentation of my work setting up the project files for the cluster and HTCondor infrastructure, as straightforward as possible, as detailed as needed.

31 The LSV Wiki's guide on HTCondor and the README files in /sopro25 as well as /sopro25/submit-files, which were written by Marius Mosbach, provide some more insight in how to use HTCondor.

32 Make sure to check that out if you need more information.

```
33
34     In case you have any questions or need help setting up HTCondor yourself, do
not hesitate to contact me, I will help you with anything you need as well as I can.
35
36
37 Some tips in advance:
38
39     - While working with the terminal / command line, make extensive use of the
autocomplete function, by pressing Tab whenever possible (unsure if Windows terminal
also has this).
40     It will autocomplete fully or partially (if in doubt), when you start
typing the name of a file or directory you want to interact with and then press Tab.
41     It will make your life much easier.
42
43     - If at any point you need to edit a file once it is already on a cluster node,
you can use the command 'nano [filename]' to edit it from the terminal (must be in the
same directory).
44
45
46
47 Some essential Linux console commands you will need:
48
49
50 pwd                -          prints out the current working directory and path
(where you are)
51
52 cd [directory]      -          go to the specified directory (you navigate, but
forward only)
53
54 cd ..              -          navigate to the parent folder ("Go back 1 step". You
can also go back more steps by using ' cd ../../ ' or ' cd ../../.. ' etc.)
55
56 cd ~                -          navigate to your home directory (' /nethome/
[yourUsername] '), from anywhere.
57
58 ls                 -          list all contents (subdirectories and files) of the
directory you are currently in
59
60 mkdir [X]           -          create a new directory X in the directory you are
currently inside of. ("Make new directory [X] here".)
61
62 rm [X]              -          delete a file or directory. Add flag ' -r ' to also
recursively delete the contents of a directory. ("Remove [X]").)
63
64 nano [file]         -          edit a file inside the terminal using the nano editor,
when you need to edit files that are already on the remote node. (nano > vim)
65
66 ssh [destination]   -          connect to some remote node via SecureShell, the LSV
cluster, for example: ' ssh [yourUsername]@contact.lsv.uni-saarland.de '
67
68 man [command]       -          displays a manual page for the specified command, where
you can read more on what it does and how to use it
69
```

70 `scp -r [sourcepath] [destination:path]` - securely copy directories /  
files via SecureShell (ssh).

71  
72 A practical example of me copying a folder (directory) from my device to the  
LSV cluster node would look like this:

73  
74 `scp -r ~/Desktop/sopro_cluster_setup/projects/sopro25`  
`awehner@contact.lsv.uni-saarland.de:/nethome/awehner/projects`

75  
76 | | |  
syntax: [command] [path on my machine to the file I want to copy] [my LSV user  
account, '@', hostname of the target, ':', destination path]

77  
78  
79  
80 Outline of the file registry structure on the LSV cluster:

81  
82 (Your account and data in the registries outlined below are available on every node.  
There are far more directories on each node, but only the following matter to us):

83  
84  
85 /  
86 |\_\_\_\_\_ /data  
87 |\_\_\_\_\_ /users  
88 |\_\_\_\_\_ /[yourUsername]  
89 |\_\_\_\_\_ /logs  
90 |\_\_\_\_\_ /sopro25  
91  
92 |\_\_\_\_\_ /nethome  
93 |\_\_\_\_\_ /[yourUsername]  
94 |\_\_\_\_\_ /projects  
95 |\_\_\_\_\_ /sopro25 (we will copy this  
directory here ourselves, details below)

96  
97  
98 This is what the file registry structure should look like. In case something does not  
work right away, please make sure all paths outlined here exist (if not, please create  
the relevant missing registries, details on how to do that below).

99  
100 Your home registry (' /~ ') will be /nethome/[yourUsername], so we will put all your  
project files in the /projects directory there (create the directory if it's not there,  
details below). In this guide, I will go into detail about how to get code running and  
submit jobs to the cluster using HTCondor (the system that manages, queues and executes  
jobs on the cluster).

101  
102  
103  
104 1. Download the sopro25 project folder from here: [LINK HERE](#)

105  
106  
107  
108 2. Open the following three files in an editor of your choice:

109

```
110         - run.sh          (in folder 'scripts')
111         - setup.sh         (in folder 'scripts')
112         - run.sub          (in folder 'submit-files')
113
114
115
116 3. In all three files, you will need to edit any occurrence of 'awehner' (my LSV
    username) with your own LSV username, to correctly set up the paths to work on your LSV
    account on the cluster.
117     here's where to edit in detail:
118
119         - run.sh           line 4
120         - setup.sh         lines 4,7,8
121         - run.sub          lines 5,6,9,10,11
122
123
124
125 4. Connect via ssh to contact.lsv.uni-saarland.de, using the following command:
126
127     ssh [yourUsername]@contact.lsv.uni-saarland.de
128
129     It will prompt you for your password upon connecting, just enter the password
    and press enter.
130
131
132
133 5. Check if the file registry on the remote node matches the outline in the graph, if
    not, create all missing directories and see that each path exists.
134
135     You will start out in /nethome/[yourUsername] , so to get to /data you have to
    navigate back to / first, using:
136
137         cd ../../ (go back 2 directories)
138
139     Then use this to get to /data/users/[yourUsername]:
140
141         cd data/users/[yourUsername]
142
143     To go back to /nethome/[yourUsername] , navigate to / using cd ../../../../ (to go
    back 3 directories) again and then use:
144
145         cd nethome/[yourUsername]
146
147     or you can use
148
149         cd ~ (the "Go back home" - command)
150
151     from anywhere to get directly back to /nethome/[yourUsername] (your home
    directory).
152
153
154
155     First navigate to /nethome[yourUsername] and check if /projects exists with
    the
```

ls command, if not, create it using the command:

```
156
157         mkdir projects
```

158 Then, use ls again to confirm you have created it there.

```
160
161 Next navigate to /data/users/[yourUsername] and check if /logs exists, if not,
create /logs with:
```

```
162
163         mkdir logs
```

164 You can again check with the ls command if it's now there.

```
166
167
168
169 6. Copy over the data from your machine to the cluster node.
```

```
170
171 Open a fresh terminal (that's not logged in to the cluster node via ssh) and copy over
the sopro25 project folder from your machine to /nethome/[yourUsername]/projects on the
remote node, using the scp -r command:
```

```
172
173 For example, from my machine it would be this (first argument: source, second
argument: destination, -r flag to recursively copy the folder's content):
```

```
174
175         scp -r ~/Desktop/sopro_cluster_setup/projects/sopro25
aweohner@contact.lsv.uni-saarland.de:/nethome/aweohner/projects
```

```
176
177 You will be prompted to enter your password again, so just enter it and confirm
with Enter.
```

```
178
179 Once you have finished copying the files, switch back to the other terminal
that's connected to the node and navigate to /nethome/[yourUsername]/projects and use
the ls command to check if the directory is there.
```

```
180
181
182 7. Build and push the docker image.
```

```
183
184 Now you need to connect (from the contact.lsv.uni-saarland.de node) to
workstation 71 of the LSV cluster, which we can use to build docker images and push
them to the LSV docker registry.
```

```
185
186 Use the command:
```

```
187
188         ssh [yourUsername]@ws71lx.lsv.uni-saarland.de
```

```
189
190 to connect to the workstation, enter your password when prompted and confirm
with enter.
```

```
191
192 Here, you can also access your data like you would on the contact node.
```

```
193
194 Navigate to /projects/sopro25/docker and build the docker image using the
following command:
```

```
195
```

```
196         IMPORTANT NOTES:
197         - You have to edit the command to include your username towards the end where I
put the [yourUsername] placeholder.
198         - When executing the command, include the whitespace and dot at the very end,
it is part of the command.
199
200         docker build -f Dockerfile --build-arg USER_UID=$UID --build-arg
USER_NAME=$(id -un) -t docker.lsv.uni-saarland.de/[yourUsername]/
dockerimage_sopro25:v1 .
201
202         Lastly, after you finished building the image, push it to the LSV docker
registry with this command (again, you will need to replace the placeholder with your
username towards the end):
203
204         docker push docker.lsv.uni-saarland.de/[yourUsername]/
dockerimage_sopro25:v1
205
206         After successfully pushing the image to the registry, log out of the
workstation node by simply typing the command 'logout' into the terminal.
207
208         You should now be back to the contact node.
209
210     8. Submit a job
211
212         From the contact node, connect to the submit node by using the following
command:
213
214         ssh [yourUsername]@submit.lsv.uni-saarland.de
215
216         Enter your password when prompted, confirm with Enter.
217
218         I have prepared a HTCondor submitfile (' run.sub ') that should work out of the
box with the paths and struture we have created so far.
219
220         Submit the job by using the following command (remember to put in your username
instead of the placeholder):
221
222         condor_submit /nethome/[yourUsername]/projects/sopro25/submit-files/
run.sub
223
224         If it does not immediately throw you any errors, it will first tell you which
cluster it has submitted your job to (a five digit number).
225         Remember this number, as it tells you which of the logfiles correspond to that
exact job submission.
226
227         You can then monitor your job with the command:
228
229         watch condor_q [yourUsername]
230
231         There you can see the status of all your jobs and check if they are still
running, on hold, rejected or finished.
232
233         Wait to see when the job finishes. It usually takes up to a minute to
complete,
```

depending on availability of the compute node we are running on.

234  
235       Once the job finishes, exit this monitor by pressing CTRL + C once, to  
interrupt the monitor process.

236  
237  
238

239 9. Check if logfiles were produced

240

241       When the job is finished running, navigate back to / and from there to data/  
users/[yourUsername]/logs/sopro25/logfiles and check if there is any files in that  
directory.

242       If everything worked correctly, there should be three files (one .log file,  
one .out file and one .err file), each named something like this:

243

244               run.sh.(XXXXX).2025\_month\_day\_somemorennumbers.log/.out/.err

245

246       The first number in parentheses (XXXXX) will be the cluster/job number from  
before, the rest will be today's date and then some more numbers.

247

248       You can check these out using the cat command, like this for example:

249

250               cat run.sh.39139.2025\_07\_28\_1753713314.out

251

252       (This was one of my output files, 39139 is the job number, the rest is the date  
and probably timestamps.)

253

254       (Also, remember to make use of the autocomplete function of your terminal by  
pressing Tab a lot.)

255

256       If you start typing 'cat run.sh' and then press Tab, it will autocomplete using  
whats in the folder until a point at which it is unsure of what file exactly you mean.

257       The first time will be until the file extension, as you now have three files  
with the same name but different file extensions in your directory.

258       Press Tab and you just have to type the extension, to access either file you  
want to see.

259       If you do this again later, when you have some older and more recent logfiles  
in the directory, it will autocomplete up to some point in the job number.

260       You will then have to type in the job number to specify, and you can hit Tab  
again so that you don't have to type the whole thing out.

261

262

263

264 10. Read all three logfiles

265

266       If all paths are working correctly as intended, each job submission will put  
out three files:

267               - a log file, containing information about the job on the cluster  
268               - an error file, which in case something goes wrong during the job will  
contain info about why it failed (file empty means the job finished as intended)

269               - an output file, which will contain the output of your actual project,  
the code you are trying to run

270

271 I have put a small example task in the project files (you can check out the  
code in /sopro25/src/mnist.py), to make the cluster actually run something.

272

273 In the output file, you should see some output at the very bottom (way below  
the long list of packages and CUDA info).

274 It should be some loss and accuracy metrics of two tiny neural networks I have  
trained and evaluated in the code.

275

276 If you can see that, the cluster setup works.

277

278

279

280 Troubleshooting in case of errors:

281

282 If something goes wrong at any point, it is most likely a path issue.

283

284 Make sure all relevant paths used in the files exist in your LSV account  
registry, or create them if necessary.

285

286 If you encounter any errors you don't know how to fix or need help in general,  
send me a message and I will do my best to help you get things running.

287

288

289

290 Whats next?

291

292

293 - Managing dependencies

294

295 Installing new dependencies requires building and pushing a new dockerimage.

296 In the Dockerfile, lines 32 and below contain the dependencies that are  
currently in the image we're using.

297 Just add whatever you need here in the same format as the ones already there,  
and docker will install and build them into the image.

298 In the case that you choose to use a new image name, the submission file  
'run.sub' needs to be updated with the correct docker image name (line 5, parameter  
'docker\_image')

299

300

301 - Getting the actual project code into the cluster

302

303 Whatever code we want to run on the cluster needs to be put in the /sopro25/src  
directory .

304 Following that, the file 'run.sh' (in /sopro25/scripts) will need to be updated  
with the new path and file name of the code to specify what code to run.

305 (line 14, under #run code; currently this specifies my example task 'mnist.py')

306

307

308 - Adapting the submitfile to request code execution on similar hardware as the  
one outlined in the Averitec shared task documents

309

310 At the moment, the submitfile is set up in a way that requests a specific  
cluster compute node and some specific hardware specs.



311           We will need to change this to more closely match the hardware outlined in the  
312           shared task (or something comparable).  
313  
314  
315   Lastly, depending on your willingness to deal with this system or not, I can always  
316   take on the job of running code and managing dependencies for the project, as I have  
317   spent some time getting into it now and have set up everything on my machine and on my  
318   cluster account.  
319  
320   This guide should just serve as a springboard into the workflow of HTCondor and the LSV  
321   cluster for you, if you want to get into it yourself.  
322  
323   If you don't want the hassle, just tell me when to pull your code from the repo and run  
324   it on the cluster, if you want me to test it for you.  
325  
326   I have uploaded the files for this on a separate GitHub repository for now (link  
327   above), so that you can clone/pull or download it more easily to test it, but I will  
328   integrate it into the GitLab repo soon(ish).  
329  
330   Also, if you have the time and nerves to spare, you should check out the other README  
331   files in the sopro25 project folder. They come from the LSV HTCondor documentation,  
332   written by Marius Mosbach, which goes a bit more into detail regarding HTCondor itself,  
333   what commands to use and how to handle it in use.  
334  
335   If you find any errors in this document or something does not work as expected, please  
336   let me know as soon as possible.  
337  
338   Alex  
339  
340   PS: Also please do me the favour of keeping a separate list of all dependencies you  
341   use, so that I can keep the docker image up to date more easily. Thanks!