

1 How to get LSV clusters running - Step by Step Guide

2

3 by Alexander Wehner

4

5 Prerequisites:

6 - OS: any Linux distribution, or alternatively WSL (Windows Subsystem for Linux) if you're currently on Windows (how to WSL: <https://learn.microsoft.com/en-us/windows/wsl/install>)

7 (ssh should also be working out of the box on the Windows command line, but I don't know if there are any intricacies with that and if it will work just as good as on Linux)

8

9 - very basic knowledge of how to use the Linux command line (necessary to connect to (using ssh) and work on the cluster nodes, which run on Ubuntu)

10 (I will walk you through the basics as good as I can)

11

12 - a copy of the project folder for the cluster that I put here (in this repository you are currently looking at):

13

14 https://github.com/alexanderwehner/SoPro_Neural_Networks_25

15

16 - LSV account, working and password changed from master to your own password

17

18 - ssh installed (should come preinstalled on pretty much any Linux disto, but if not, install via the command 'apt install ssh' (e.g. for Ubuntu)

19 or use any other package manager of your choice)

20

21 - note: wherever I have put the [yourUsername] placeholder, replace it with just your username, drop the square brackets (and don't put any whitespaces in or something like that)

22

23

24 Disclaimer:

25

26 I have explicitly written this guide assuming you don't know how to operate a Linux terminal, to make sure everyone understands what we are doing, even without having any Linux experience.

27 I tried to go into great detail on each step, explaining the commands you need and how to use them, even for the most basic commands.

28 If you already are using Linux and have some experience working with the Linux terminal, much of this guide will feel redundant and overly detailed.

29 I made this decision because I wanted to make sure that everyone can not just reproduce but also understand what they are doing, even if they have no Linux experience.

30

31 This guide is the compiled documentation of my work setting up the project files for the cluster and HTCondor infrastructure, as straightforward as possible, as detailed as needed.

32 The LSV Wiki's guide on HTCondor and the README files in /sopro25 as well as /sopro25/submit-files, which were written by Marius Mosbach, provide some more insight in how to use HTCondor.

33 Make sure to check that out if you need more information.

34

35 In case you have any questions or need help setting up HTCondor yourself, do
36 not hesitate to contact me, I will help you with anything you need as well as I can.
37
38 Some tips in advance:
39
40 - While working with the terminal / command line, make extensive use of the
 autocomplete function, by pressing Tab whenever possible (unsure if Windows terminal
 also has this).
41 It will autocomplete fully or partially (if in doubt), when you start
 typing the name of a file or directory you want to interact with and then press Tab.
42 It will make your life much easier.
43
44 - If at any point you need to edit a file once it is already on a cluster node,
 you can use the command 'nano [filename]' to edit it from the terminal (must be in the
 same directory).
45
46
47
48 Some essential Linux console commands you will need:
49
50
51 pwd - prints out the current working directory and path
 (where you are)
52
53 cd [directory] - go to the specified directory (you navigate, but
 forward only)
54
55 cd .. - navigate to the parent folder ("Go back 1 step". You
 can also go back more steps by using ' cd ../../ ' or ' cd ../../.. ' etc.)
56
57 cd ~ - navigate to your home directory (' /nethome/
 [yourUsername] '), from anywhere.
58
59 ls - list all contents (subdirectories and files) of the
 directory you are currently in
60
61 mkdir [X] - create a new directory X in the directory you are
 currently inside of. ("Make new directory [X] here".)
62
63 rm [X] - delete a file or directory. Add flag ' -r ' to also
 recursively delete the contents of a directory. ("Remove [X]").
64
65 nano [file] - edit a file inside the terminal using the nano editor,
 when you need to edit files that are already on the remote node. (nano > vim)
66
67 ssh [destination] - connect to some remote node via SecureShell, the LSV
 cluster, for example: ' ssh [yourUsername]@contact.lsv.uni-saarland.de '
68
69 man [command] - displays a manual page for the specified command, where
 you can read more on what it does and how to use it
70
71 scp -r [sourcepath] [destination:path] - securely copy directories /

files via SecureShell (ssh).

A practical example of me copying a folder (directory) from my device to the LSV cluster node would look like this:

```
scp -r ~/Desktop/sopro_cluster_setup/projects/sopro25  
aweohner@contact.lsv.uni-saarland.de:/nethome/aweohner/projects
```

```
|                                     |                                     |  
syntax:      [command] [path on my machine to the file I want to copy] [my LSV user  
account, '@', hostname of the target, ':', destination path]
```

Outline of the file registry structure on the LSV cluster:

(Your account and data in the registries outlined below are available on every node. There are far more directories on each node, but only the following matter to us):

```
/
|_____ /data
|           |_____ /users
|                   |_____ /[yourUsername]
|                               |_____ /logs
|                                       |_____ /sopro25
|_____ /nethome
|           |_____ /[yourUsername]
|                   |_____ /projects
|                               |_____ /sopro25 (we will copy this
directory here ourselves, details below)
```

This is what the file registry structure should look like. In case something does not work right away, please make sure all paths outlined here exist (if not, please create the relevant missing registries, details on how to do that below).

Your home registry (' /~ ') will be /nethome/[yourUsername], so we will put all your project files in the /projects directory there (create the directory if it's not there, details below). In this guide, I will go into detail about how to get code running and submit jobs to the cluster using HTCondor (the system that manages, queues and executes jobs on the cluster).

1. Download the sopro25 project folder from here: https://github.com/alexanderwehner/SoPro_Neural_Networks_25

2. Open the following three files in an editor of your choice:

```
111         - run.sh          (in folder 'scripts')
112         - setup.sh        (in folder 'scripts')
113         - run.sub         (in folder 'submit-files')
114
115
116
117 3. In all three files, you will need to edit any occurrence of 'aweher' (my LSV
username) with your own LSV username, to correctly set up the paths to work on your LSV
account on the cluster.
118     here's where to edit in detail:
119
120         - run.sh          line 4
121         - setup.sh        lines 4,7,8
122         - run.sub         lines 5,6,9,10,11
123
124
125
126 4. Connect via ssh to contact.lsv.uni-saarland.de, using the following command:
127
128     ssh [yourUsername]@contact.lsv.uni-saarland.de
129
130     It will prompt you for your password upon connecting, just enter the password
and press enter.
131
132
133
134 5. Check if the file registry on the remote node matches the outline in the graph, if
not, create all missing directories and see that each path exists.
135
136     You will start out in /nethome/[yourUsername] , so to get to /data you have to
navigate back to / first, using:
137
138         cd ../../ (go back 2 directories)
139
140     Then use this to get to /data/users/[yourUsername]:
141
142         cd data/users/[yourUsername]
143
144     To go back to /nethome/[yourUsername] , navigate to / using cd ../../.. (to go
back 3 directories) again and then use:
145
146         cd nethome/[yourUsername]
147
148     or you can use
149
150         cd ~ (the "Go back home" - command)
151
152     from anywhere to get directly back to /nethome/[yourUsername] (your home
directory).
153
154
155
156     First navigate to /nethome[yourUsername] and check if /projects exists with
the
```

```
157     ls command, if not, create it using the command:
158         mkdir projects
159
160     Then, use ls again to confirm you have created it there.
161
162     Next navigate to /data/users/[yourUsername] and check if /logs exists, if not,
163     create /logs with:
164         mkdir logs
165
166     You can again check with the ls command if it's now there.
167
168
169
170 6. Copy over the data from your machine to the cluster node.
171
172 Open a fresh terminal (that's not logged in to the cluster node via ssh) and copy over
173 the sopro25 project folder from your machine to /nethome/[yourUsername]/projects on the
174 remote node, using the scp -r command:
175
176     For example, from my machine it would be this (first argument: source, second
177     argument: destination, -r flag to recursively copy the folder's content):
178         scp -r ~/Desktop/sopro_cluster_setup/projects/sopro25
179     awehner@contact.lsv.uni-saarland.de:/nethome/awehner/projects
180
181     You will be prompted to enter your password again, so just enter it and confirm
182     with Enter.
183
184     Once you have finished copying the files, switch back to the other terminal
185     that's connected to the node and navigate to /nethome/[yourUsername]/projects and use
186     the ls command to check if the directory is there.
187
188
189
190 7. Build and push the docker image.
191
192 Now you need to connect (from the contact.lsv.uni-saarland.de node) to
193 workstation 71 of the LSV cluster, which we can use to build docker images and push
194 them to the LSV docker registry.
195
196     Use the command:
197         ssh [yourUsername]@ws71lx.lsv.uni-saarland.de
198
199     to connect to the workstation, enter your password when prompted and confirm
200     with enter.
201
202     Here, you can also access your data like you would on the contact node.
203
204     Navigate to /projects/sopro25/docker and build the docker image using the
205     following command:
```

197 IMPORTANT NOTES:

198 - You have to edit the command to include your username towards the end where I
put the [yourUsername] placeholder.

199 - When executing the command, include the whitespace and dot at the very end,
it is part of the command.

200

201 docker build -f Dockerfile --build-arg USER_UID=\$UID --build-arg
USER_NAME=\$(id -un) -t docker.lsv.uni-saarland.de/[yourUsername]/
dockerimage_sopro25:v1 .

202

203 Lastly, after you finished building the image, push it to the LSV docker
registry with this command (again, you will need to replace the placeholder with your
username towards the end):

204

205 docker push docker.lsv.uni-saarland.de/[yourUsername]/
dockerimage_sopro25:v1

206

207 After successfully pushing the image to the registry, log out of the
workstation node by simply typing the command 'logout' into the terminal.

208

209 You should now be back to the contact node.

210

211 8. Submit a job

212

213 From the contact node, connect to the submit node by using the following
command:

214

215 ssh [yourUsername]@submit.lsv.uni-saarland.de

216

217 Enter your password when prompted, confirm with Enter.

218

219 I have prepared a HTCondor submitfile (' run.sub ') that should work out of the
box with the paths and struture we have created so far.

220

221 Submit the job by using the following command (remember to put in your username
instead of the placeholder):

222

223 condor_submit /nethome/[yourUsername]/projects/sopro25/submit-files/
run.sub

224

225 If it does not immediately throw you any errors, it will first tell you which
cluster it has submitted your job to (a five digit number).

226 Remember this number, as it tells you which of the logfiles correspond to that
exact job submission.

227

228 You can then monitor your job with the command:

229

230 watch condor_q [yourUsername]

231

232 There you can see the status of all your jobs and check if they are still
running, on hold, rejected or finished.

233

234 Wait to see when the job finishes. It usually takes up to a minute to
complete,

depending on availability of the compute node we are running on.

235

236 Once the job finishes, exit this monitor by pressing CTRL + C once, to
interrupt the monitor process.

237

238

239

240 9. Check if logfiles were produced

241

242 When the job is finished running, navigate back to / and from there to data/
users/[yourUsername]/logs/sopro25/logfiles and check if there is any files in that
directory.

243 If everything worked correctly, there should be three files (one .log file,
one .out file and one .err file), each named something like this:

244

245 run.sh.(XXXXX).2025_month_day_somemorennumbers.log/.out/.err

246

247 The first number in parentheses (XXXXX) will be the cluster/job number from
before, the rest will be today's date and then some more numbers.

248

249 You can check these out using the cat command, like this for example:

250

251 cat run.sh.39139.2025_07_28_1753713314.out

252

253 (This was one of my output files, 39139 is the job number, the rest is the date
and probably timestamps.)

254

255 (Also, remember to make use of the autocomplete function of your terminal by
pressing Tab a lot.)

256

257 If you start typing 'cat run.sh' and then press Tab, it will autocomplete using
whats in the folder until a point at which it is unsure of what file exactly you mean.

258 The first time will be until the file extension, as you now have three files
with the same name but different file extensions in your directory.

259 Press Tab and you just have to type the extension, to access either file you
want to see.

260 If you do this again later, when you have some older and more recent logfiles
in the directory, it will autocomplete up to some point in the job number.

261 You will then have to type in the job number to specify, and you can hit Tab
again so that you don't have to type the whole thing out.

262

263

264

265 10. Read all three logfiles

266

267 If all paths are working correctly as intended, each job submission will put
out three files:

268 - a log file, containing information about the job on the cluster

269 - an error file, which in case something goes wrong during the job will
contain info about why it failed (file empty means the job finished as intended)

270 - an output file, which will contain the output of your actual project,
the code you are trying to run

271

272 I have put a small example task in the project files (you can check out the
code in /sopro25/src/mnist.py), to make the cluster actually run something.

273

274 In the output file, you should see some output at the very bottom (way below
the long list of packages and CUDA info).

275 It should be some loss and accuracy metrics of two tiny neural networks I have
trained and evaluated in the code.

276

277 If you can see that, the cluster setup works.

278

279

280

281 Troubleshooting in case of errors:

282

283 If something goes wrong at any point, it is most likely a path issue.

284

285 Make sure all relevant paths used in the files exist in your LSV account
registry, or create them if necessary.

286

287 If you encounter any errors you don't know how to fix or need help in general,
send me a message and I will do my best to help you get things running.

288

289

290

291 Whats next?

292

293

294 - Managing dependencies

295

296 Installing new dependencies (the way I have set it up here) requires building
and pushing a new dockerimage each time you update the dependency list in the
dockerfile.

297 Just update the Dockerfile and repeat step 7.

298 In the Dockerfile, lines 32 and below contain the dependencies that are
currently in the image we're using.

299 Just add whatever you need here in the same format as the ones already there,
and docker will install and build them into the image.

300 In the case that you choose to use a new image name, the submission file
'run.sub' needs to be updated with the correct docker image name (line 5, parameter
'docker_image')

301

302

303 - Getting the actual project code into the cluster

304

305 Whatever code we want to run on the cluster needs to be put in the /sopro25/src
directory .

306 Following that, the file 'run.sh' (in /sopro25/scripts) will need to be updated
with the new path and file name of the code to specify what code to run.

307 (line 14, under #run code; currently this specifies my example task 'mnist.py')

308

309

310 - Adapting the submitfile to request code execution on specific hardware specs.

311

312 At the moment, the submitfile is set up in a way that requests a specific
cluster compute node and some specific hardware specs.

313 Lines 12-15 of the submitfile will need to be updated to specify what exact
hardware specs are needed for the job at hand.

314

315

316

317 This guide should just serve as a springboard into the workflow of HTCondor and the LSV
cluster for you, if you want to get into it yourself.

318

319 Also, if you have the time and nerves to spare, you should check out the other README
files in the sopro25 project folder. They come from the LSV HTCondor documentation,
written by Marius Mosbach, which goes a bit more into detail regarding HTCondor itself,
what commands to use and how to handle it in use.

320

321 The project files I used and modified for my own setup are also based on Mr. Mosbach's
HTCondor documentation (currently the official working LSV cluster documentation).

322

323 If you find any errors in this document or something does not work as expected, please
let me know as soon as possible.

324

325 Alex